

# Scala & Spark ETL — Complete Study Material (Beginner → Advanced)

**Goal:** Teach you how to design, implement, test, and operate robust ETL pipelines using Scala + Apache Spark. Includes beginner-friendly explanations, code snippets, illustrations, and recommended GitHub repositories to clone and study.

---

## 1. Table of contents

1. Overview & Learning Path
  2. Prerequisites & Tooling
  3. Getting Started — Hands-on Setup
  4. Minimal ETL: CSV → Parquet → PostgreSQL (step-by-step)
  5. Project structure & best practices
  6. Testing, CI, and local development
  7. Advanced topics: Partitioning, Bucketing, Schema Evolution, Performance
  8. Production concerns: Orchestration, Observability, Security
  9. Recommended GitHub repositories (beginner → advanced)
  10. Mini projects & exercises
  11. Appendix: Useful sbt, Docker, SQL snippets
- 

## 2. Overview & Learning Path

- **Phase 1 (Basics):** Scala syntax, sbt, small Spark jobs, read/write CSV/Parquet.
- **Phase 2 (Intermediate):** Transformations, joins, window functions, CTE-like logic using DataFrame APIs, UDFs/UDAFs.
- **Phase 3 (Advanced):** Optimize jobs (partitioning, broadcast joins), streaming (Structured Streaming), orchestration (Airflow), production practices (logging, metrics, retries).

Estimated time: 4–8 weeks of focused learning (self-paced).

---

## 3. Prerequisites & Tooling

**Languages & frameworks:** Scala (2.12.x or 2.13.x), Apache Spark 3.x

**Tools:**

- sbt (build tool)
- Git
- PostgreSQL (or any RDBMS) for sinks
- Docker (optional, for uniform environment)
- IDE: IntelliJ IDEA (Community) + Scala plugin
- Local Spark: start Spark with `spark-submit` or use `sbt run` for small apps

**Recommended IDE settings:**

- Enable Scala plugin, set SDK to Java 11 or 17 depending on Spark build.
- 

## 4. Getting Started — Hands-on Setup

### Install tools (brief)

- Install **Java 11** or 17.
- Install **sbt**: follow sbt docs or package manager.
- Install **PostgreSQL** (or use SQLite for simple tests).
- Install **Git** and clone repositories.

### Sample project scaffold (sbt)

`build.sbt` minimal:

```
name := "simple-etl"
```

```
version := "0.1.0"
```

```
scalaVersion := "2.12.17"
```

```
libraryDependencies ++= Seq(  
  "org.apache.spark" %% "spark-core" % "3.4.1",  
  "org.apache.spark" %% "spark-sql" % "3.4.1",  
  "org.postgresql" % "postgresql" % "42.6.0"  
)
```

```
assembly / mainClass := Some("com.example.etl.Main")
```

Adjust Spark and Scala versions if you use a corporate cluster. Keep Spark 3.x for modern compatibility.

## Project layout

```
simple-etl/
├── src/main/scala/com/example/etl/
│   ├── Main.scala
│   ├── config/ConfigReader.scala
│   ├── jobs/CsvToParquetJob.scala
│   └── utils/SparkSessionBuilder.scala
├── src/test/scala/
├── build.sbt
└── conf/application.conf
```

---

## 5. Minimal ETL Example: CSV → Parquet → PostgreSQL

This example shows a readable, runnable template you can clone and expand.

### SparkSession builder (utils/SparkSessionBuilder.scala)

```
package com.example.etl.utils
```

```
import org.apache.spark.sql.SparkSession
```

```
object SparkSessionBuilder {
  def build(appName: String = "simple-etl") = {
    SparkSession.builder()
      .appName(appName)
      .master("local[*]")
      .config("spark.sql.shuffle.partitions", "4")
      .getOrCreate()
  }
}
```

### Main entry (Main.scala)

```
package com.example.etl
```

```
import com.example.etl.jobs.CsvToParquetJob
import com.example.etl.utils.SparkSessionBuilder
```

```
object Main {
  def main(args: Array[String]): Unit = {
```

```

val spark = SparkSessionBuilder.build("csv-to-parquet")
CsvToParquetJob.run(spark, args)
spark.stop()
}
}

```

## **Job: read CSV, transform, write Parquet and Postgres (jobs/CsvToParquetJob.scala)**

```
package com.example.etl.jobs
```

```
import org.apache.spark.sql.{DataFrame, SparkSession}
import org.apache.spark.sql.functions._
```

```
object CsvToParquetJob {
```

```

def run(spark: SparkSession, args: Array[String]): Unit = {
val input = if (args.nonEmpty) args(0) else "data/input/invoices.csv"
val parquetOut = if (args.length > 1) args(1) else "data/output/invoices.parquet"

```

```

val df = readCsv(spark, input)
  .withColumn("invoice_amount", col("quantity") * col("unit_price"))
  .withColumn("invoice_date", to_date(col("invoice_date"), "yyyy-MM-dd"))

```

```
df.write.mode("overwrite").parquet(parquetOut)
```

```

writeToPostgres(df, "jdbc:postgresql://localhost:5432/accounting", "invoices")
}

```

```

def readCsv(spark: SparkSession, path: String): DataFrame = {
spark.read
  .option("header", "true")
  .option("inferSchema", "true")
  .csv(path)
}

```

```

def writeToPostgres(df: DataFrame, jdbcUrl: String, table: String): Unit = {
val props = new java.util.Properties()
props.setProperty("user", "postgres")
props.setProperty("password", "postgres")

```

```

df.write
  .mode("append")
  .jdbc(jdbcUrl, table, props)
}

```

```
}
```

### How to run locally

1. Put a sample CSV at `data/input/invoices.csv` (small sample included below).
2. `sbt run` or `sbt "run data/input/invoices.csv data/output/invoices.parquet"`

### Sample CSV (invoices.csv)

```
invoice_id,customer_id,invoice_date,quantity,unit_price
1,100,2024-07-01,2,150.00
2,101,2024-07-03,1,200.00
3,100,2024-07-05,3,50.00
```

---

## 6. Project Structure & Best Practices

- **Separation of concerns:** Keep `config`, `io`, `transform`, `jobs` directories.
- **Idempotency:** Jobs should be re-runnable (use *write modes*, staging areas, or check markers).
- **Config-driven:** Use `application.conf` or environment variables for DB creds and paths.
- **Logging:** Use `log4j` or `slf4j`. Avoid `println` for production.
- **Small Tasks:** Design small, testable transformations.
- **Use DataFrame/Dataset API** instead of RDDs for declarative optimizations.

### Example config (application.conf)

```
app {
  env = "local"
  input.path = "data/input"
  output.path = "data/output"
  jdbc.url = "jdbc:postgresql://localhost:5432/accounting"
}
```

---

## 7. Testing & CI

- Unit tests using `scalatest` or `munit`.
- Use `spark-testing-base` or create `SparkSession` with local master for tests.

### Sample test

```
import org.scalatest.funsuite.AnyFunSuite
import com.example.etl.utils.SparkSessionBuilder

class CsvToParquetJobTest extends AnyFunSuite {
  val spark = SparkSessionBuilder.build("test")

  test("invoice_amount is calculated") {
    import spark.implicits._
    val df = Seq((1,2,100.0)).toDF("invoice_id","quantity","unit_price")
    val withAmount = df.withColumn("invoice_amount", col("quantity") * col("unit_price"))
    assert(withAmount.collect().head.getAs[Double]("invoice_amount") == 200.0)
  }
}
```

#### CI tips:

- Run `sbt test` in CI (GitHub Actions). Use matrix for Java/Scala versions.
  - Optionally build Docker image with `sbt assembly` and `spark-submit` for integration tests.
- 

## 8. Advanced Topics & Performance

### Partitioning & Bucketing

- Partition by `date` (e.g., `year=`, `month=`) for large fact tables.
- Use bucketing for join-heavy datasets with fixed key cardinality.

#### Write partitioned parquet

```
df.write.partitionBy("year", "month").parquet("/data/facts/invoices")
```

### Broadcast joins

Use `broadcast(df)` for small dimension tables.

```
import org.apache.spark.sql.functions.broadcast
val joined = fact.join(broadcast(dim), Seq("id"))
```

### Caching & Persisting

Cache intermediate DataFrames when reused multiple times:

```
val cached = df.filter(...).cache()
cached.count() // materialize
```

## File formats

- Use Parquet for columnar efficiency.
- Consider **Delta Lake** or **Apache Hudi** for ACID + schema evolution.

## Monitoring & Metrics

- Push job metrics to Prometheus (via Spark metrics or JVM exporter).
  - Emit custom counters using  
`org.apache.spark.SparkContext.longAccumulator`.
- 

# 9. Production Concerns: Orchestration, Observability, Security

## Orchestration

- Use **Apache Airflow** (Python) to schedule and orchestrate Spark jobs. Provide a sample DAG to call `spark-submit` or submit jobs to a cluster.

### Sample Airflow DAG (concept)

```
from airflow import DAG
from airflow.operators.bash_operator import BashOperator
from datetime import datetime
```

```
with DAG('etl_pipeline', start_date=datetime(2025,1,1), schedule_interval='@daily') as dag:
    run_etl = BashOperator(
        task_id='run_spark_job',
        bash_command='spark-submit --class com.example.etl.Main /opt/etl/simple-etl.jar'
    )
```

## Observability

- Centralized logs (ELK/EFK).
- Track job duration, input row counts, error counts.
- Alert when job fails or input counts deviate from baseline.

## Security

- Use secrets manager for DB credentials (Vault, AWS Secrets Manager).
  - Limit JDBC user privileges.
  - TLS between services where possible.
- 

## 10. Recommended GitHub Repositories (Beginner → Advanced)

Clone these to study patterns and sample pipelines.

### Beginner / Learning

- [SETL \(Scala ETL framework\)](https://github.com/SETL-Framework/setl) — [github.com/SETL-Framework/setl](https://github.com/SETL-Framework/setl)
  - Why: Clean, modular; good to learn structure.
- [etl-spark](https://github.com/alexland/etl-spark) — [github.com/alexland/etl-spark](https://github.com/alexland/etl-spark)
  - Why: Minimalistic, good for a first real pipeline.
- [MyDataFramework](https://github.com/vbounyasit/MyDataFramework) — [github.com/vbounyasit/MyDataFramework](https://github.com/vbounyasit/MyDataFramework)
  - Why: Framework-oriented reusable pipelines.

### Intermediate

- [spark-etl-framework](https://github.com/qwshen/spark-etl-framework) — [github.com/qwshen/spark-etl-framework](https://github.com/qwshen/spark-etl-framework)
  - Why: Focused on end-to-end ingestion → transformation.
- [spark-etl](https://github.com/aphp/spark-etl) — [github.com/aphp/spark-etl](https://github.com/aphp/spark-etl)
  - Why: Integrates Spark with Postgres/DB sinks.

### Advanced / Production

- [Teams-League-Airflow-Spark-Scala-ETL](https://github.com/tosun-si/teams-league-airflow-spark-scala-etl) — [github.com/tosun-si/teams-league-airflow-spark-scala-etl](https://github.com/tosun-si/teams-league-airflow-spark-scala-etl)
    - Why: Shows orchestration + cloud storage + BigQuery patterns.
  - [Scala-and-Spark-in-Practice](https://github.com/ruslanmv/Scala-and-Spark-in-Practice) — [github.com/ruslanmv/Scala-and-Spark-in-Practice](https://github.com/ruslanmv/Scala-and-Spark-in-Practice)
    - Why: Exercises and performance patterns.
- 

## 11. Mini Projects & Exercises

1. **Invoice ETL** — CSV invoices → clean → Parquet partitioned by year/month → load to Postgres. Add tests.



2. **Revenue Dashboard Feed** — Aggregate monthly revenue and write a materialized table.
  3. **Duplicate Detector** — Find and report duplicate invoices (same customer, amount, date).
  4. **Streaming variant** — Use Structured Streaming to ingest new transactions and update a running ledger.
  5. **Data Quality Framework** — Create checks (row counts, null ratio, unique key checks) and store results.
- 

## 12. Appendix: Useful Commands & Snippets

### SBT assembly

sbt clean assembly

### Build fat JAR and run with spark-submit

```
spark-submit --class com.example.etl.Main --master local[*]  
target/scala-2.12/simple-etl-assembly-0.1.0.jar data/input/invoices.csv  
data/output/invoices.parquet
```

\*\*Export to CSV from Postgre