

Spark Optimization Techniques



⋮

Code-Level Optimizations

a. Avoid Using `collect()` Unnecessarily

❌ Bad
data = df.collect()

✅ Good
df.show(10) # for preview
df.take(10) # for small samples

Code-Level Optimizations

b. Use mapPartitions Instead of map

```
# ❌ map (inefficient for DB connections)  
rdd.map(lambda x: db_write(x))
```

```
# ✅ mapPartitions  
def write_partition(iter):  
    conn = create_db_connection()  
    results = [db_write(record, conn) for record in iter]  
    conn.close()  
    return results  
  
rdd.mapPartitions(write_partition)
```

Code-Level Optimizations

c. Use broadcast for Small Lookup Tables

```
from pyspark.sql import functions as F
```

```
# ❌ Normal join (causes shuffle)  
df_large.join(df_small, "id")
```

```
# ✅ Broadcast join  
from pyspark.sql import SparkSession  
from pyspark.sql.functions import broadcast
```

```
df_large.join(broadcast(df_small), "id")
```

Code-Level Optimizations

d. Cache/Persist Reused Data

```
# ❌ Without cache → recomputes each time  
df_filtered = df.filter(df["age"] > 30)  
print(df_filtered.count())  
print(df_filtered.collect())
```

```
# ✅ With cache  
df_filtered = df.filter(df["age"] > 30).cache()
```

Code-Level Optimizations

e. Avoid UDFs, Use Spark SQL Functions

```
from pyspark.sql.functions import col, upper  
  
# ❌ UDF  
from pyspark.sql.functions import udf  
from pyspark.sql.types import StringType  
@udf(returnType=StringType())  
def to_upper(s): return s.upper()  
df.withColumn("name_upper", to_upper(col("name")))  
  
# ✅ Built-in function  
df.withColumn("name_upper", upper(col("name")))
```

Code-Level Optimizations

f. Partition Data Before Writing

```
# ❌ Single output file (bad parallelism)  
df.write.csv("output.csv")
```

```
# ✅ Partition by column (parallel, scalable)  
df.write.partitionBy("year",  
"month").parquet("/output/path/")
```

Code-Level Optimizations

g. Optimize Shuffle Partitions

```
# Default 200 partitions → too many small files  
spark.conf.set("spark.sql.shuffle.partitions", "200")
```

```
# ✅ Tune based on data size  
spark.conf.set("spark.sql.shuffle.partitions", "50")
```

Code-Level Optimizations

h. Push Down Filters / Column Pruning

❌ Select all columns
df = spark.read.parquet("data.parquet")

✅ Select only needed
df = spark.read.parquet("data.parquet").select("id",
"name").filter("age > 30")



Code-Level Optimizations

i. Repartition vs Coalesce

```
df.repartition(100) # Increase partitions  
df.coalesce(10)   # Reduce partitions
```



Cluster-Level Optimizations

⋮

Cluster-Level Optimizations

a. Executor & Core Tuning

- Don't overload executors with too many cores.
- Rule of thumb: 5 cores per executor.

--executor-memory 4G

--executor-cores 5

--num-executors 10



Cluster-Level Optimizations

b. Use Dynamic Allocation

- Enable Spark to scale resources automatically.

```
--conf spark.dynamicAllocation.enabled=true  
--conf spark.shuffle.service.enabled=true  
--conf spark.dynamicAllocation.minExecutors=2  
--conf spark.dynamicAllocation.maxExecutors=50
```

Cluster-Level Optimizations

c. Choose Right File Format

- Prefer Parquet/ORC over CSV/JSON (columnar, compressed).
- Use Snappy compression for Parquet (splittable & fast).

```
df.write.mode("overwrite").parquet("/path/output/",  
compression="snappy")
```

Cluster-Level Optimizations

d. Data Locality

- Co-locate Spark with HDFS nodes.
- Ensure `spark.locality.wait` is tuned (default = 3s).

`--conf spark.locality.wait=1s`

Cluster-Level Optimizations

e. Adaptive Query Execution (AQE)

- Introduced in Spark 3.x – optimizes joins/shuffles at runtime.

```
--conf spark.sql.adaptive.enabled=true
```

```
--conf spark.sql.adaptive.shuffle.targetPostShuffleInputSize=64MB
```

Cluster-Level Optimizations

g. Serialization

- Use Kryo instead of Java serialization (faster, smaller).

```
--conf spark.serializer=org.apache.spark.serializer.KryoSerializer
```

Cluster-Level Optimizations

h. Shuffle & Spill Optimization

- Store shuffle data on SSDs if possible.
- Increase memory fraction for shuffle:

--conf spark.memory.fraction=0.6

Summary Code-Level

01

Avoid collect(), prefer actions (take, show).

02

Use mapPartitions, broadcast, cache, SQL functions.

03

Filter early, select fewer columns.

04

Optimize shuffle partitions, repartition/coalesce wisely.

•
•
•



Cluster-Level Optimizations

h. Shuffle & Spill Optimization

- Store shuffle data on SSDs if possible.
- Increase memory fraction for shuffle:

--conf spark.memory.fraction=0.6

Summary Cluster-Level

01

Tune executor memory/cores.

02

Enable dynamic allocation.

03

Use Parquet/ORC + Snappy. Use Kryo serialization.

04

Enable AQE and Handle small files & shuffle efficiently.

⋮

