```
spark-submit \
--class com.example.App \
--master yarn \
--deploy-mode cluster \
--num-executors 30 \
--executor-cores 5 \
--executor-memory 19g \
--conf spark.yarn.executor.memoryOverhead=2048 \
app.jar
```

Default Partition
128 MB Size

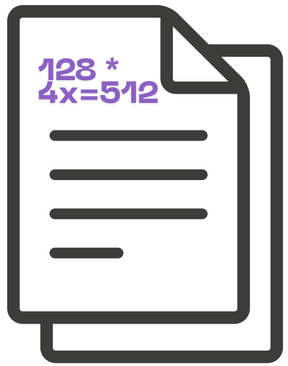

128 MB Partition Size

Default Partition
128 MB

10 GB File = 10240 MB *
128 MB
80 Partitions

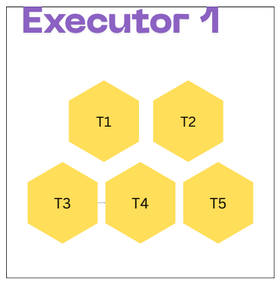

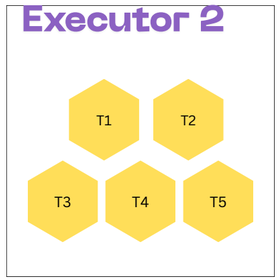

10240
MB

Assign a min of 4X memory for each core.
128 * 4 = 512 MB

Max Allowed CPU Cores
for each Executors

No of Executors = Total Core / Executor Cores

Multiply it by executor cores to get executor memory
512 MB * 5 Cores
= 2560 MB [3 GB]

512 MB X 5 Cores = 2560 MB = 3 GB

Partition 1 | Partition 2 | ..... | Partition 80